

# Parallelism – Some Theory

Computational Science II (CAAM 520)

---

Christopher Thiele

Rice University, Spring 2020

# Overview

- Why parallelize?
- Shared and distributed memory parallelism
- Processes and threads
- OpenMP and MPI
- Terminology
- Strong and weak scalability
- Amdahl's law

# Why parallelize?

There are (at least) two possible motivations:

- Utilization of more computational resources to reduce time-to-solution.
- Utilization of more computational resources to solve larger problems.

Furthermore, contemporary hardware, e.g., multi-core CPUs, makes parallelization somewhat mandatory.

# Shared memory parallelism

Shared memory parallelism targets platforms on which all compute units have access to the same memory.

Examples:

- Multiple CPU cores on a single chip accessing main memory
- Multiple CPUs who can access each other's memory
- GPGPUs

# Distributed memory parallelism

Distributed memory parallelism targets platforms on which each compute unit has its own, private memory.

Example:

- A computer cluster in which compute nodes are connected via network.

# Shared and distributed memory parallelism

The boundaries are fluid due to technologies such as ***remote direct memory access*** (RDMA).

Besides the hardware perspective, shared and distributed memory parallelism can also be viewed as software concepts.

→ Just because compute units can share memory, they do not have to.

# Processes and threads

Modern operating systems (OS) for desktop and laptop computers, servers, etc. support multi-tasking, i.e., tasks or applications can run concurrently.

Multi-tasking can be implemented in different ways:

- Time-sharing: All tasks run on the same CPU (core), and the OS switches between them.
- True multi-tasking: Tasks run on separate CPUs or CPU cores, i.e., truly concurrently.

# Processes and threads

**Processes** are independent tasks that can be scheduled by the OS using either multi-tasking approach.

**Threads** are tasks that are created by a process. They typically and share a process ID, memory, etc. with the parent process, but they can be scheduled independently.

→ Specific differences between processes and threads depend on the operating system.



# Processes and threads

Threads are suitable for shared memory environments.

- We will use OpenMP for shared memory parallelism.
- OpenMP applications spawn multiple threads which access shared memory.

Distributed memory environments require processes.

- We will use the Message Passing Interface (MPI) to develop distributed memory parallel applications.
- MPI applications spawn multiple processes which exchange messages.

# Terminology

Some conventions:

- **Thread**: the software concept discussed before, i.e., not a hardware feature as in hyperthreading
- **CPU** or **processor**: an entire physical CPU, may consist of multiple CPU cores
- We always distinguish a **process** (software) from a **processor** (hardware).
- **(Compute) node**: an entire computer, i.e., the whole box

# Terminology

Consider a parallel application which does the (normalized) amount of work

$$s + p = 1,$$

where

- $p$  is the fraction of work that can be performed in ***parallel***, and
- $s$  is the fraction of work that must be performed ***sequentially***.

# Terminology

Let  $T_f^1$  be the time required to do the **fixed** amount of work  $s + p$  **sequentially**.

Let  $T_f^N = s$  be the time required to do the same amount of work  $s + p$  in **parallel** using  $N$  **workers**.

(**Note:** Our definitions differ from those used by Hager & Wellein.)

# Terminology

For  $N \geq 1$ , let us define **performance** as work over time, i.e.,

$$P_f^N = \frac{s + p}{T_f^N}.$$

# Terminology

We define the **speedup** as

$$S^N = \frac{P_f^N}{P_f^1} = \frac{T_f^1}{T_f^N}.$$

→ The speedup tells us how much performance increases, and how much time-to-solution decreases when using  $N$  workers instead of one.

# Terminology

We define ***parallel efficiency*** as

$$E_f^N = \frac{P_f^N}{NP_f^1} = \frac{T_f^1}{NT_f^N} = \frac{S^N}{N}.$$

When using  $N$  workers instead of one, we should expect a speedup of at most  $N$ .

→ Parallel efficiency tells us what fraction of the ideal speedup can be achieved.

# Strong scalability

So far we kept the amount of work fixed, and we increased the number of workers.

The speedup, as defined for this setting, measures ***strong scalability***.



# Weak scalability

Let us now increase the number of workers, keeping the amount of work per worker fixed.

In other words, consider the time  $T_v^N$  it takes to do the **variable** amount of work

$$N = N \underbrace{(s + p)}_{=1}$$

using  $N$  workers.

(Note that  $T_v^1 = T_f^1$  by definition.)

# Weak scalability

Let us define performance and efficiency as

$$P_v^N = \frac{N}{T_v^N}$$

and

$$E_v^N = \frac{P_v^N}{P_v^1} = \frac{T_v^1}{T_v^N}.$$

The efficiency, as defined for a variable amount of work, measures **weak scalability**.

→ Compare it to the definition of  $E_f^N = \frac{T_f^1}{NT_f^N}$  in the strong scalability setting!

# Amdahl's law

Let us separate  $T_f^N$  into

$$T_f^N = T_{f,s} + T_{f,p}^N,$$

where

- $T_{f,s}$  is the time required to do the fraction  $s$  of the total amount of work  $s + p = 1$  which cannot be parallelized, and
- $T_{f,p}^N$  is the time required to do the remaining fraction  $p$  of the work in parallel.

# Amdahl's law

Assume perfect<sup>1</sup> strong scalability, i.e.,

$$T_{f,p}^N = \frac{T_{f,p}^1}{N}.$$

Then the speedup is

$$S^N = \frac{T_f^1}{T_f^N} = \frac{T_f^1}{T_{f,s} + T_{f,p}^N} = \frac{T_f^1}{T_{f,s} + \frac{T_{f,p}^1}{N}} = \frac{T_f^1}{T_{f,s} + \frac{T_f^1 - T_{f,s}}{N}}.$$

---

<sup>1</sup>In some circumstances, **superlinear** scaling can be observed, i.e.,  $T_{f,p}^N < T_{f,p}^1/N$ .

# Amdahl's law

After simplifying, we obtain

$$S^N = \frac{T_f^1}{T_{f,s} + \frac{T_f^1 - T_{f,s}}{N}} = \frac{T_f^1}{\left(1 - \frac{1}{N}\right) T_{f,s} + \frac{T_f^1}{N}} = \frac{1}{\left(1 - \frac{1}{N}\right) s + \frac{1}{N}},$$

i.e.,

$$S^N = \frac{1}{s + \frac{1-s}{N}}$$

(Amdahl's law).

# Amdahl's law

Amdahl's law tells us the speedup assuming perfect strong scalability.

In particular, it implies that the speedup is limited, because

$$S^N = \frac{1}{s + \frac{1-s}{N}} \xrightarrow{N \rightarrow \infty} \frac{1}{s}.$$

# Amdahl's law

**Example:** Suppose that 90% of an application's work can be done in parallel, while 10% ( $s = 0.1$ ) must be done sequentially.

Assuming perfect strong scalability, the speedup achieved when using  $N = 10$  workers instead of one is

$$S^N = \frac{1}{s + \frac{1-s}{N}} = \frac{1}{0.1 + \frac{0.9}{10}} \approx 5.26.$$

Furthermore, the speedup can never exceed

$$\frac{1}{s} = \frac{1}{0.1} = 10.$$

# Amdahl's law

More generally, Amdahl's law says that

$$S = \frac{1}{s + \frac{1-s}{S_p}},$$

where  $S_p$  is the speedup experienced by the fraction  $p$  of the total amount of work.

So far we assumed  $S_p = N$ .



# Amdahl's law

In its general form, Amdahl's law is useful in many contexts.

**Example:** By code optimization, 43% of the total work in a sequential code can be accelerated by a factor of 1.6.

Then the speedup for the overall application is

$$S = \frac{1}{0.57 + \frac{0.43}{1.6}} \approx 1.19.$$

# Amdahl's law

In its general form, Amdahl's law is useful in many contexts.

**Example:** An application spends 17% of its time reading from and writing to a hard drive.

The hard drive is replaced with a newer drive that is faster by a factor of 3.2.

Then the speedup for the overall application is

$$S = \frac{1}{0.83 + \frac{0.17}{3.2}} \approx 1.13.$$